

# Gaming the Two Numbers They Control: Quantization, Weight-Cut Bunching, and the Limits of Strength in 3.9 Million Powerlifting Results

Adir Elmakais   David Levin

Statistical Theory, M.Sc. in Computer Science, Bar-Ilan University

**Abstract**—A powerlifting competitor controls exactly two numbers: the weight loaded on the bar and their bodyweight at weigh-in. We ask whether these choices leave measurable statistical fingerprints, using  $\sim 3.94$  million competition results from the public-domain OpenPowerlifting database. Four findings, each answered with a course-appropriate test and led by effect size (at  $n \approx 3.9\text{M}$  almost everything is “significant”). (i) *Quantization*: 96.2% of 13.4M attempt loads sit exactly on the 2.5 kg plate grid; a generalized likelihood-ratio test (GLRT) whose null parameter lies on the boundary is calibrated by a parametric bootstrap. (ii) *Weight-cut bunching*: bodyweight piles up just below every IPF class limit (de-heaped  $\log(\text{below/above}) = +1.92$  at 83 kg); a Cattaneo–Jansson–Ma density-discontinuity test rejects at all seven limits (surviving Holm and Benjamini–Hochberg), a non-limit control shows no bunching, and placebos show no false positives. (iii) *Allometry*: strength scales with bodyweight differently by sex ( $b = 0.75$  men,  $0.51$  women vs. the isometric  $2/3$ ; a  $\text{Sex} \times \log(\text{BW})$  interaction is highly significant and persists on common support). (iv) *Prediction*: basic controllable variables predict the total ( $R^2 = 0.70$ , random forest, leakage-safe grouped CV). These choices leave measurable traces in the data. Code and data: <https://github.com/adirelm/opl-stat-theory>.

## I. INTRODUCTION

Powerlifting comprises three lifts (squat, bench press, deadlift), each with three attempts, contested within bodyweight classes. Two features make the sport an unusually clean laboratory for strategic behaviour. First, the implements are *discrete*: plates come in fixed denominations, so the load a competitor can place on the bar is not a continuous choice. Second, eligibility is governed by a *threshold*: a lifter who weighs even slightly above a class limit is bumped into a heavier, stronger field. A competitor’s agency therefore reduces to two numbers: the load chosen for each attempt, and the bodyweight at weigh-in. Both interact with a sharp administrative rule. We ask whether these two choices leave detectable statistical signatures, and what  $\sim 3.94\text{M}$  results reveal about the limits of human strength.

**Related work.** Manipulation of a “running variable” around an administrative threshold produces *bunching*, a standard tool in public economics [1] that is classically tested with the McCrary density-discontinuity test [2] and its modern, bias-corrected local-polynomial form due to Cattaneo, Jansson, and Ma [3]. Rapid pre-weigh-in weight cutting is well documented in weight-class sports. On the physiological side, the square-cube law predicts that strength, governed by muscle cross-sectional area, scales with bodyweight as  $\propto \text{BW}^{2/3}$ . Extreme-

value theory has been used to estimate the ultimate limits of athletic performance from record data [6]. On OpenPowerlifting specifically, recent work studies *scoring and scaling* systems [7]; to our knowledge, no formal manipulation test of weight-cut bunching has been applied to this dataset.

**Contribution.** The project follows a single thread, “the two numbers and the limits of strength,” and uses both halves of the course, classical inference and machine learning. Our contributions are: (1) a rounded-likelihood *quantization* GLRT for the plate grid, with a parametric-bootstrap calibration that respects the boundary nature of the null; (2) to our knowledge the first formal density-discontinuity test of weight-cut *bunching* on OpenPowerlifting, hardened with a non-limit control, placebo cutoffs, de-heaping, and multiple-testing corrections; (3) a sex-interaction *allometry* model that tests the scaling exponent against the isometric  $2/3$ ; and (4) a leakage-safe strength-*prediction* model. Throughout we lead with effect sizes and confidence intervals rather than  $p$ -values, as explained next.

**Methodological stance.** At  $n \approx 3.9\text{M}$  the power to detect any non-zero effect is effectively one, so a  $p$ -value answers a question we do not care about (“is the effect exactly zero?”) rather than the one we do (“is it large?”). We therefore report effect sizes with 95% confidence intervals as the primary evidence, use  $p$ -values only as a secondary screen, and apply all four standard multiple-testing corrections (Bonferroni, Holm, and Šidák for familywise error, and Benjamini–Hochberg for the false-discovery rate) to every pre-declared family of tests. Because the same lifter appears in many rows, naive standard errors understate uncertainty (pseudo-replication); the formal tests therefore use one row per lifter or cluster-robust standard errors, declared per hypothesis.

## II. RESULTS

All intervals are 95%. Headline effect sizes use all rows; formal tests use one row per lifter (to respect independence) and, where the class scheme matters, the modern kg-only era (2014+).

### A. Attempt loads are quantized to the 2.5 kg grid (H1)

Of 13,397,702 valid attempt loads, **96.20%** (CI [96.19, 96.21]) fall exactly on the 2.5 kg plate grid (Fig. 1); the smallest standard plate is 1.25 kg, loaded on both sides, so 2.5 kg is the minimum increment. To test this formally we restrict to the kg-native (0.5 kg) lattice and model the




Fig. 1. H1: attempt loads concentrate on the 2.5 kg plate grid.

within-cell position as a mixture that places weight  $\pi$  on the grid point and  $1 - \pi$  on a uniform background. The GLRT of  $\pi = 0$  versus  $\pi > 0$  has its null parameter on the boundary of  $[0, 1]$ , so Wilks’ theorem fails: the null distribution of  $2 \ln \Lambda$  is the 50:50 mixture  $\frac{1}{2}\chi_0^2 + \frac{1}{2}\chi_1^2$  of Chernoff and Self–Liang [4], [5]. Our parametric bootstrap recovers exactly this (point mass 0.51 at zero; calibrated 95% cutoff 2.59 versus the naive 3.84). The observed statistic,  $2 \ln \Lambda = 3.9 \times 10^7$ , exceeds every bootstrap draw ( $p < 3.3 \times 10^{-4}$ ), and a  $\chi^2$  goodness-of-fit against a uniform within-cell distribution is overwhelming ( $5.0 \times 10^7$ ,  $df = 4$ ). The 3.8% off-grid loads are largely a units artifact: of the 2.1% that fall off the 0.5 kg lattice entirely, 82% are round pound loads (e.g., 102.06 kg = 225 lb), themselves quantized on a different grid.

### B. Bodyweight bunches below class limits (H2)

Bodyweight piles up just below real class limits (Fig. 2). On de-heaped data (exact round/half-kilogram weigh-ins removed, so the signal is not mere digit preference), the log-ratio of just-below to just-above counts is positive at all seven IPF men’s limits and strongest at 83 kg ( $+1.92 \pm 0.04$ , a  $6.8\times$  asymmetry; Table I, Fig. 3); a non-limit control at 91 kg is flat-to-slightly-negative ( $-0.21$ ). The formal Cattaneo–Jansson–Ma density-discontinuity test [3] (one random meet per lifter, 2014+) rejects at all seven limits with the density dropping at the cutoff ( $t < 0$ ), surviving both Holm and Benjamini–Hochberg correction; the control shows no excess below (a small *opposite*-direction discontinuity), and no placebo shows the bunching pattern. The effect is robust across equipment (raw  $+2.14$ , equipped  $+1.51$ ) and tested status (tested  $+1.93$ ), and persists on the clean kg-only subset ( $+1.75$ ). The weakest bunching is at the 120 kg cutoff, which is adjacent to the unbounded 120+ class where the incentive to cut is weakest.

**Robustness and falsification.** The result survives four checks. (i) *De-heaping*: removing exact round/half-kilogram weigh-ins leaves the asymmetry intact, so it is manipulation, not digit preference. (ii) *Control*: the non-limit point at 91 kg shows no excess below. (iii) *Placebos*: cutoffs placed between real limits show no bunching-direction discontinuity after correction. (iv) *Subgroups and correction*: the effect holds within raw, equipped, and tested subsets, and the formal density test

TABLE I  
DE-HEAPED  $\log(\text{BELOW}/\text{ABOVE})$  AT EACH IPF MEN’S LIMIT (ALL ROWS).  
LAST ROW IS A NON-LIMIT CONTROL.

| Limit (kg)   | log-ratio    | $\pm$ | $\times$ asymmetry |
|--------------|--------------|-------|--------------------|
| 59           | +0.99        | 0.04  | 2.7                |
| 66           | +1.18        | 0.04  | 3.2                |
| 74           | +1.00        | 0.03  | 2.7                |
| <b>83</b>    | <b>+1.92</b> | 0.04  | <b>6.8</b>         |
| 93           | +1.65        | 0.04  | 5.2                |
| 105          | +1.17        | 0.04  | 3.2                |
| 120          | +0.74        | 0.06  | 2.1                |
| 91 (control) | −0.21        | 0.03  | 0.8                |

rejects at all seven limits even under the strictest (Bonferroni) correction.

### C. Allometric scaling differs by sex (H3)

Fitting  $\log(\text{Total}) = a + b \log(\text{BW})$  (Fig. 4), the exponent is  $b = 0.75$  (CI  $[0.74, 0.75]$ ) for men and  $0.51$  ( $[0.50, 0.51]$ ) for women, with heteroskedasticity-robust (HC3) standard errors on one row per lifter. A Wald test rejects the isometric  $b = 2/3$  for both sexes (men  $z = +46.8$ , women  $z = -64.1$ ); the one-sided framing is motivated by Neyman–Pearson/UMP logic, exact only under a one-parameter exponential-family model. A pooled  $\text{Sex} \times \log(\text{BW})$  interaction estimates the men-minus-women slope gap at  $+0.24$  ( $p \ll 10^{-3}$ ). The gap also *grows* on the common bodyweight support (0.84 men vs. 0.38 women), so it is not an artifact of the sexes occupying different weight ranges. The Pearson correlation of  $\log \text{BW}$  and  $\log \text{Total}$  is 0.60 (Fisher-transform CI  $[0.60, 0.60]$ ; Spearman 0.58). The message is not that women are weaker, but that the *estimated scaling* differs; we leave the mechanism to future work.

### D. Strength is predictable from basic variables (H4)

Predicting the total from controllable/basic variables (bodyweight, sex, equipment, age), with every split grouped by lifter and all preprocessing fit inside each training fold, a random forest reaches  $R^2 = 0.70$  (RMSE 86 kg) versus 0.61 (RMSE 98 kg) for linear regression (Fig. 5); permutation importance is dominated by sex and bodyweight, and continuous-feature VIFs are  $\approx 1$  (no collinearity concern). A logistic “made-weight” classifier built from *non-bodyweight* features (age, tested status, era, equipment; the Dots score is excluded because it is a function of total and bodyweight) reaches  $\text{AUC} = 0.58$  (PR-AUC 0.88, balanced accuracy 0.54). The signal is weak, as expected: cutting is driven by bodyweight relative to the limit, which defines the label and which these features omit.

### E. Supporting analyses

**Distribution structure.** The apparent two-component shape of raw totals is just sex: by BIC and a bootstrap likelihood-ratio test, raw Total prefers two Gaussian components ( $\Delta \text{BIC} = 1.4 \times 10^3$ ) but the sex-normalized Dots score prefers one (Fig. 6); the “mixture” is sex, and normalizing removes it. **Tested vs. untested.** Untested competitions show higher

**Fig 2. Bodyweight bunches just below a real class limit, not at a non-limit control**




Fig. 2. H2: bunching just below a real limit (83 kg) but not at a non-limit control (91 kg); round-number weigh-ins removed.




Fig. 3. H2: excess just-below mass at every real limit; control near zero.

**Fig 4. Allometric scaling by sex (all-row OLS fits; formal per-lifter HC3: 0.75 / 0.51)**




Fig. 4. H3: allometric scaling by sex vs. the isometric 2/3.

Dots than tested ones (medians 344 vs. 318, Mann-Whitney  $p \rightarrow 0$ , rank-biserial  $-0.20$ ), consistent with the expected effect of performance-enhancing drugs. **Extreme values.** A GEV fit to equal-size block maxima of recent totals gives shape  $\xi = -0.02$  with bootstrap CI  $[-0.20, 3.77]$ : a population strength *ceiling is not supported*, since the CI includes zero. **Time trend.** Median Dots drifted down over 2000–2024 (Spearman  $\rho = -0.54$ ), consistent with broadening participation rather than declining strength. **Sequential detection.**

A Wald sequential probability ratio test on the chronological stream of just-below/above indicators accepts the bunching hypothesis after a stopping time of only  $N = 27$  (Fig. 8). **Independence.** After deduplicating to one row per lifter, the association between tested status and cutting is negligible (Cramér’s  $V = 0.008$ ), illustrating how pseudo-replication inflates significance. **Residual normality.** The H3 residuals are left-skewed with a heavy lower tail (skew  $-0.77$ ; Fig. 7), so they are not exactly normal. This does not undermine the slope: at this  $n$  its sampling distribution is normal by the central limit theorem, and we report HC3-robust standard errors. A formal normality test rejects beyond the smallest samples, as expected at this scale, so we treat it as a diagnostic and lead with effect sizes.

### III. METHODS

**Data and preprocessing.** We use a pinned snapshot of OpenPowerlifting (3,941,811 rows, public domain; the SHA-256 is recorded for exact reproducibility). The unit of analysis is the *attempt* for H1 (the mechanism is within-competitor, so the grid is robust to clustering) and one *row per lifter* for the cross-sectional tests; *de-heaping* drops exact round/half-kg weigh-ins so the bunching signal is not digit preference. For prediction, all cross-validation is grouped by lifter so the same athlete never appears in both train and test.

**GLRT and the boundary (H1).** With log-likelihood  $\ell$ , the statistic is  $2 \ln \Lambda = 2(\ell_1 - \ell_0)$ . Under regularity  $2 \ln \Lambda \rightarrow \chi^2_r$ , but when the null fixes a parameter on the boundary of its space the limit becomes a mixture  $\frac{1}{2}\chi^2_0 + \frac{1}{2}\chi^2_1$  [4], [5]; we estimate the null distribution directly by simulating data under the fitted null model (parametric bootstrap) and comparing the observed statistic to the simulated quantiles. Concretely, for within-cell position  $u \in \{0, \dots, 4\}$  on the 0.5 kg lattice the model is

$$P(u) = \pi \mathbb{I}[u = 0] + (1 - \pi) \frac{1}{5}, \quad \pi \in [0, 1], \quad (1)$$

**Fig 5. Predicting strength: bodyweight + sex dominate; RF beats linear**




Fig. 5. H4: strength prediction; importance and grouped-CV  $R^2$ .

**Fig 6. The apparent 'mixture' in raw Total is just sex; Dots removes it**




Fig. 6. The apparent “mixture” in raw Total is sex; Dots removes it.

and the GLRT contrasts  $\pi = 0$  (no grid preference, on the boundary) with the constrained MLE  $\hat{\pi} \geq 0$ .

**Density discontinuity (H2).** For a cutoff  $c$  the manipulation test contrasts the one-sided densities,  $\theta = \ln \hat{f}^-(c) - \ln \hat{f}^+(c)$ , estimated by local polynomials without pre-binning [2], [3];  $\theta < 0$  ( $t < 0$ ) indicates excess mass just below  $c$ . We complement the formal test with a de-heaped log count-ratio as the interpretable effect size.

**Allometry (H3).** OLS on log-log data gives the scaling exponent  $b$ ; we report HC3/cluster-robust standard errors, a Wald statistic  $W = (\hat{b} - 2/3)/\widehat{\text{se}}(\hat{b})$ , and a pooled Sex  $\times$  log(BW) interaction whose coefficient is exactly the men-minus-women slope difference. Correlations use the Fisher  $z$ -transform for confidence intervals.

**Learning and the rest of the toolbox.** Prediction uses linear and random-forest regression and logistic classification, evaluated by grouped cross-validation ( $R^2$ , adjusted  $R^2$ , RMSE, MAE, permutation importance, VIF, AUC/PR-

AUC). Supporting analyses use a bootstrap mixture-LRT with AIC/BIC, a GEV extreme-value fit, a Wald sequential probability ratio test (with decision boundaries  $A = \ln \frac{1-\beta}{\alpha}$ ,  $B = \ln \frac{\beta}{1-\alpha}$ , a stopping time, and an expected-sample-size approximation), one- and two-sample nonparametric tests (Wilcoxon signed-rank; Mann-Whitney), an  $F$ -test/Kruskal-Wallis omnibus with post-hoc comparisons, and a  $\chi^2$  independence test with Cramér’s  $V$  and standardized residuals. The extreme-value fit uses the GEV distribution

$$G(x) = \exp \left\{ - \left( 1 + \xi \frac{x-\mu}{\sigma} \right)^{-1/\xi} \right\}, \quad (2)$$

where  $\xi < 0$  implies a bounded upper tail (a strength ceiling) and  $\xi \geq 0$  does not; we estimate  $\xi$  from equal-size block maxima and bootstrap its interval. The sequential test’s expected sample size under the alternative is approximated by  $\mathbb{E}[N \mid H_1] \approx [(1-\beta)A + \beta B] / [p_1 s_1 + (1-p_1)s_0]$ , with  $s_0, s_1$  the per-observation log-likelihood increments. Table II maps the course toolbox onto the study.

**Fig 8. H3 residuals: a heavy lower tail (non-normal); the slope is CLT-robust at large  $n$**




Fig. 7. H3 residuals: heavy lower tail; the slope is CLT-robust at large  $n$ .

**Fig 7. Sequential test (stopping time) and a-priori power**




Fig. 8. Sequential test (stopping time  $N = 27$ ) and a-priori power ( $n = 211$  for power 0.90).

**Reproducibility.** Every result is produced by a small set of scripts from the pinned snapshot; numbers are written to machine-readable files, figures are exported at 300 DPI, and an independent re-derivation from the raw data reproduced every headline value reported here.

#### IV. DISCUSSION

**Conclusions.** The two numbers a competitor controls leave measurable, falsification-surviving traces. Attempt loads are quantized to the plate grid; bodyweight bunches specifically below class limits, and not at controls or placebos in the bunching direction. Beyond strategy, the data speak to the limits of strength: scaling with bodyweight is sex-specific and departs from the isometric  $2/3$ , and basic controllable variables already explain  $\sim 70\%$  of the variance in the total. The narrative moves from description to formal test to prediction using one coherent dataset and a single story.

**Limitations and null results.** This is an observational study: we identify population-level patterns consistent with

weight cutting, not the intent of any individual. We report three null results as genuine findings. The unbounded 120+ class shows the weakest bunching, exactly as incentives predict; the extreme-value analysis does *not* support a strength ceiling (the  $\xi$  confidence interval includes zero); and we find no material association between tested status and cutting once pseudo-replication is removed. On the methods side, lifter identity is name-based, the per-lifter meet is sampled at random (we report sensitivity), and the formal bunching test is restricted to the modern kg-only era. Several extensions remain for future work: more federations and eras, a structural bunching model that recovers the share of manipulators, and a multiverse robustness check.

#### REFERENCES

- [1] H. J. Kleven, "Bunching," *Annual Review of Economics*, vol. 8, pp. 435–464, 2016.
- [2] J. McCrary, "Manipulation of the running variable in the regression discontinuity design: A density test," *J. Econometrics*, vol. 142, no. 2, pp. 698–714, 2008.

TABLE II  
COVERAGE OF THE COURSE TOOLBOX.

| Concept                                | Where used          |
|----------------------------------------|---------------------|
| GLRT, MLE, $\chi^2$ goodness-of-fit    | H1 grid model       |
| Boundary LRT, parametric bootstrap     | H1 null calibration |
| Density discontinuity (McCrary/CJM)    | H2 bunching         |
| Multiple testing (Bonf./Holm/Šidák/BH) | H2 limits, H3       |
| OLS, $R^2$ , Wald, interaction term    | H3 allometry        |
| One-/two-sided, one-sample tests       | H1, H3              |
| Pearson/Spearman, Fisher CI            | H3                  |
| Regression, random forest, logistic    | H4 prediction       |
| Learning metrics, grouped CV, VIF      | H4                  |
| Mixture LRT, AIC/BIC                   | structure           |
| Mann–Whitney (two-sample)              | tested vs. untested |
| Wilcoxon signed-rank (paired)          | attempt 1 vs. 3     |
| $F$ /ANOVA, Kruskal–Wallis, post-hoc   | equipment           |
| $\chi^2$ independence, Cramér’s $V$    | tested $\times$ cut |
| Extreme-value theory (GEV)             | strength ceiling    |
| Sequential Wald (SPRT), stopping time  | cut detection       |
| Power, Type I/II tradeoff              | a-priori power      |

- [3] M. D. Cattaneo, M. Jansson, and X. Ma, “Simple local polynomial density estimators,” *J. Amer. Statist. Assoc.*, vol. 115, no. 531, pp. 1449–1455, 2020.
- [4] H. Chernoff, “On the distribution of the likelihood ratio,” *Ann. Math. Statist.*, vol. 25, no. 3, pp. 573–578, 1954.
- [5] S. G. Self and K.-Y. Liang, “Asymptotic properties of maximum likelihood estimators and likelihood ratio tests under nonstandard conditions,” *J. Amer. Statist. Assoc.*, vol. 82, no. 398, pp. 605–610, 1987.
- [6] J. H. J. Einmahl and J. R. Magnus, “Records in athletics through extreme-value theory,” *J. Amer. Statist. Assoc.*, vol. 103, no. 484, pp. 1382–1391, 2008.
- [7] Peyen *et al.*, “Discussing diminishing returns: A new scoring system for powerlifting,” arXiv:2503.13040, 2025.
- [8] OpenPowerlifting, <https://www.openpowerlifting.org> (data released into the public domain).